

Team ID: Team Name:

Intent Classification in Banking Queries: A Comparative Study of Classical, Embedding-based, and Transformer Models

Abstract

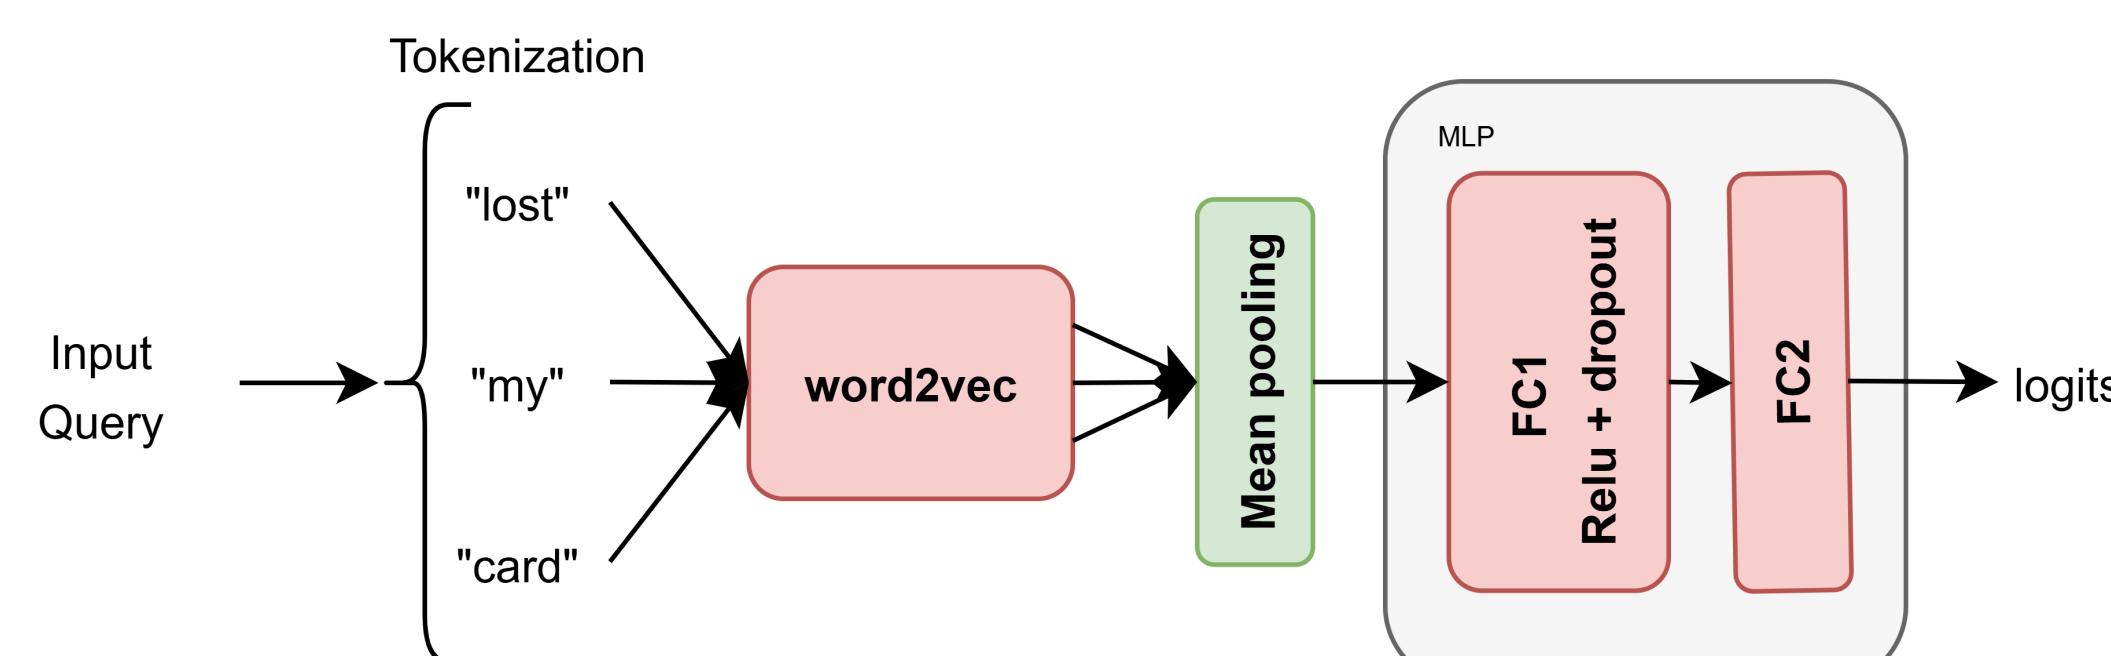
- We tackle intent classification for short banking queries using the BANKING77 dataset^[1] (77 intents).
- We compare three model families: TF/TF-IDF + Multinomial NB, Word2Vec sentence embeddings + MLP, and fine-tuned transformer encoders (BERT / DistilBERT), all trained on the same split and evaluated with macro-F1.
- Transformer models achieve the best performance: BERT with CLS pooling reaches macro-F1 ≈ 0.95 , while TF-based and Word2Vec baselines remain competitive (macro-F1 ≈ 0.86 and 0.88) at substantially lower computational cost.

Goal of the Project

- Build an intent classifier that maps short banking queries to 77 predefined BANKING77 intent classes.
- Compare classical, embedding-based, and transformer models on BANKING77 and analyze their results to understand performance trade-offs.

Methodology

- Dataset**
 - BANKING77**: 13k short English banking queries labeled into 77 intent classes.
 - Clean, curated text: no noise \rightarrow only light preprocessing needed.
 - Imbalanced dataset \rightarrow we report macro-F1, macro-precision, and macro-recall.
- Preprocessing (per method)**
 - Method 1**: default preprocessing from CountVectorizer / TfIdfVectorizer (lowercasing, tokenization, n-grams).
 - Method 2**: Lowercasing, quote normalization, and regex-based word tokenization; sentence = mean of Word2Vec embeddings (OOV \rightarrow UNK).
 - Method 3**: Raw text passed to the pretrained tokenizer for subword tokenization and truncation (max length 64), with dynamic padding at the batch level.
- Data split**: 70% train / 15% validation / 15% test, using a fixed random seed and the same split.
- Method 1 TF / TF-IDF + Multinomial Naive Bayes**
 - Approach**
 - Convert queries to sparse bag-of-words vectors (TF / TF-IDF).
 - Train a Multinomial Naive Bayes classifier using grid search.
 - Evaluate the best two models on the test set.
 - Tuning**
 - 5-fold stratified grid search.
 - Searched n -grams, \min_df , \max_df , sublinear TF (TF-IDF), and NB smoothing α .
 - Objective: macro-F1 (class-imbalance aware).
 - Best Configurations**
 - TF: unigrams + bigrams, $\min_df = 2$, $\alpha = 0.1$.
 - TF-IDF: unigrams, $\min_df = 3$, sublinear TF = True, $\alpha = 0.1$.
- Method 2 Word2Vec Sentence Embeddings + MLP**
 - Approach**
 - Pretrained word2vec-google-news-300.
 - Sentence embedding = mean of word vectors.
 - Train a small MLP classifier (ReLU + dropout).
 - Tuning**
 - Hyperparameter tuning on the validation macro-F1.
 - Best configuration** : Hidden=512, Dropout=0.1, Batch=64, lr= 0.001, weight_decay=0, label_smoothing=0.1
 - OOV Handling**
 - Variant 1: Include OOV tokens.
 - Variant 2: Ignore OOV tokens.
- Method 3 Fine-Tuned Transformer Encoders (BERT / DistilBERT)**
 - Approach**
 - Fine-tune two pretrained encoders: BERT-base-uncased and DistilBERT-base-uncased.
 - Sentence representation (pooling strategy):
 - CLS token.
 - Mean pooling.
 - Apply a Linear layer as classifier head.
 - Tuning**
 - Jointly fine-tune encoder + classification head using AdamW and cross-entropy.
 - Hyperparameter tuning on validation macro-F1.
 - Best configuration**:
 - BERT (CLS / AVG): lr = 3e-5, weight_decay = 0.0, warmup_ratio = 0.06, batch_size = 64, epochs = 15.
 - DistilBERT (CLS): lr = 3e-5, weight_decay = 0.01, warmup_ratio = 0.06, batch_size = 64, epochs = 15.



- Tuning**
 - Hyperparameter tuning on the validation macro-F1.
 - Best configuration** : Hidden=512, Dropout=0.1, Batch=64, lr= 0.001, weight_decay=0, label_smoothing=0.1
- OOV Handling**
 - Variant 1: Include OOV tokens.
 - Variant 2: Ignore OOV tokens.

Method 3 Fine-Tuned Transformer Encoders (BERT / DistilBERT)

- Approach**
 - Fine-tune two pretrained encoders: BERT-base-uncased and DistilBERT-base-uncased.
 - Sentence representation (pooling strategy):
 - CLS token.
 - Mean pooling.
 - Apply a Linear layer as classifier head.
- Tuning**
 - Jointly fine-tune encoder + classification head using AdamW and cross-entropy.
 - Hyperparameter tuning on validation macro-F1.
 - Best configuration**:
 - BERT (CLS / AVG): lr = 3e-5, weight_decay = 0.0, warmup_ratio = 0.06, batch_size = 64, epochs = 15.
 - DistilBERT (CLS): lr = 3e-5, weight_decay = 0.01, warmup_ratio = 0.06, batch_size = 64, epochs = 15.

Results

Metrics

- Metrics reported on the test split: macro-recall, macro-precision, macro-F1 (averaged over 77 intents).

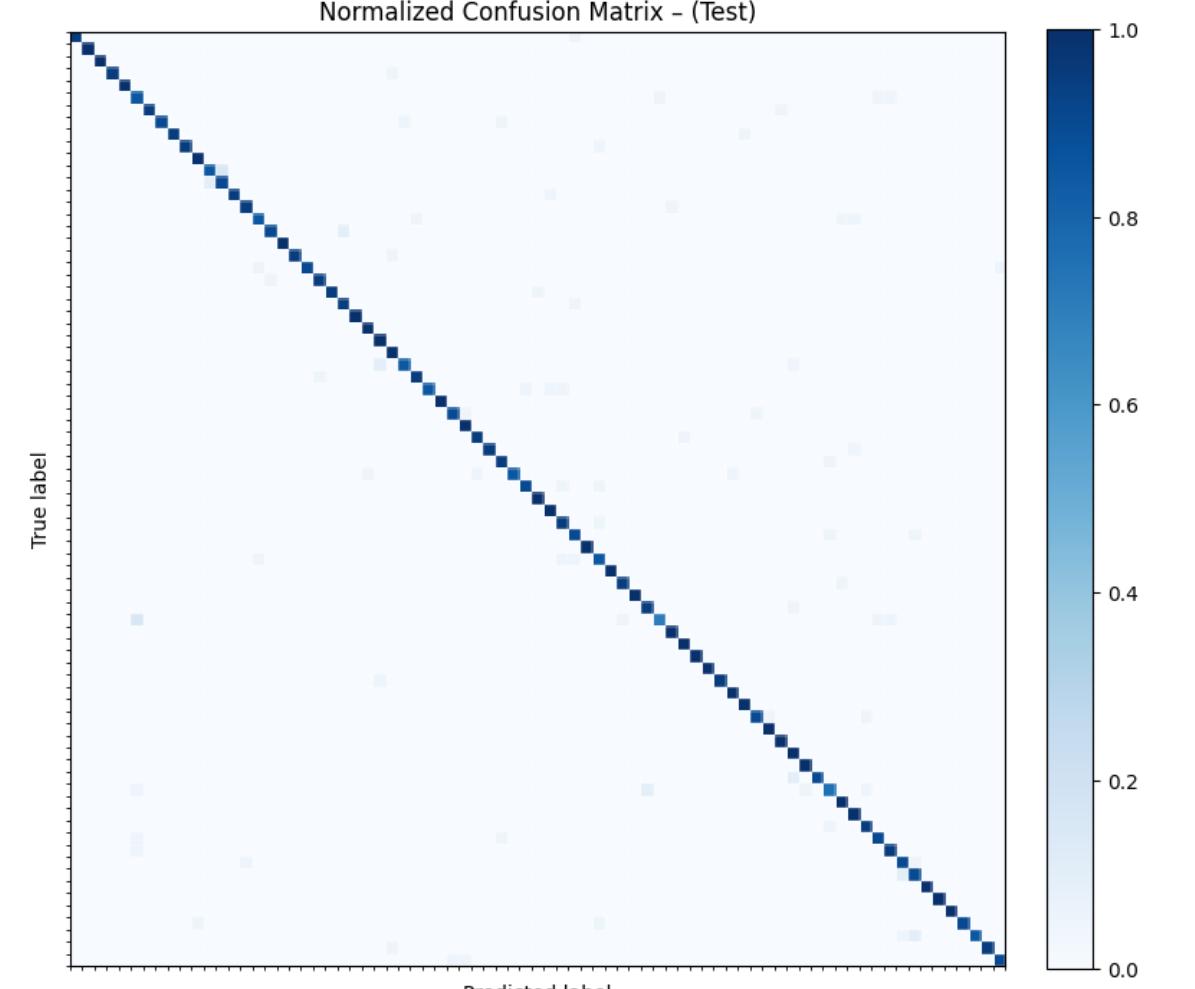
Overall performance

Method	Model	Recall	Precision	F1-score
Baselines	BERT [2]	0.926	0.931	0.926
	distilbert [3]	-	-	0.924
	RoBERTa [4]	0.935	0.937	0.935
1	TF + Multinomial NB	0.865	0.871	0.865
	TF-IDF + Multinomial NB	0.855	0.864	0.855
2	Word2Vec + MLP (with oov)	0.875	0.886	0.875
	Word2Vec + MLP (ignoring oov)	0.879	0.889	0.879
	DistilBERT (cls)	0.930	0.933	0.930
3	BERT (cls)	0.945	0.947	0.945
	BERT (avg)	0.940	0.943	0.940

Confusion Matrix

- Strong diagonal dominance indicates highly accurate class separation.
- Very low off-diagonal density \rightarrow misclassifications are rare and scattered.
- No major overlapping classes, reflecting stable generalization across all 77 intents.

Fig: Normalized confusion matrix for the best model.
Normalized Confusion Matrix - (Test)



Key Observations

1. First Method:

- TF performs slightly better than TF-IDF on BANKING77.
- Best TF model reaches macro-F1 0.865, a strong classical baseline.

2. Second Method:

- Word2Vec sentence embeddings + MLP improve performance over bag-of-words.
- Ignoring OOV tokens yields a small but consistent gain (0.875 \rightarrow 0.879 macro-F1).

3. Third Method:

- All transformer variants clearly outperform Methods 1 and 2.
- DistilBERT (CLS) reaches macro-F1 = 0.930, BERT (avg) reaches 0.940, and BERT (CLS) reaches 0.945, with BERT clearly surpassing DistilBERT and CLS pooling slightly ahead of mean pooling.

4. Overall:

- Performance improves from sparse (TF) \rightarrow dense static embeddings (Word2Vec) \rightarrow contextual transformers.
- Method 1 gives a strong, efficient baseline; Method 2 adds semantic structure at moderate cost; Method 3 achieves the best performance, matching and exceeding published transformer baselines.

Analysis

First Method:

- TF performs slightly better than TF-IDF, which may be due to short, repetitive queries where down-weighting frequent terms removes useful signal.
- Best TF setup (unigrams + bigrams, $\min_df = 2$, $\alpha = 0.1$) may be due to bigrams capturing common patterns and \min_df removing noisy rare n-grams.
- Overall performance (≈ 0.86 macro-F1) is strong for a classical model, showing it handles this dataset surprisingly well.

Second Method:

- Word2Vec improves over TF/TF-IDF, which may be due to dense semantic embeddings reducing sparsity and giving similar embeddings to related words, unlike TF/TF-IDF's sparse counts with no semantic structure.
- Ignoring OOV tokens slightly outperforms using a UNK embedding, which may be due to avoiding distortion in the averaged sentence representation.
- Both variants outperform Method 1, and the best result (≈ 0.879 macro-F1) is solid for such a lightweight model, which is also expected given the added semantic information.

Third Method:

- Transformers outperform Methods 1 and 2, explained by the fact that they learn contextual embeddings using multi-head attention, unlike TF/TF-IDF's sparse counts or Word2Vec's static vectors.
- BERT surpasses DistilBERT, explained by its larger depth and representational capacity, leading to stronger feature extraction.
- Within BERT, CLS pooling slightly beats mean pooling: the CLS token is explicitly trained as a sequence summary, while averaging can dilute important tokens.
- We outperform the published Banking77 baselines, which may be due to more careful fine-tuning (hyperparameter tuning, longer training, and using the best pooling choice, CLS).

Compute trade-off: TF/TF-IDF (fastest) \rightarrow Word2Vec \rightarrow DistilBERT \rightarrow BERT (slowest, most accurate), as expected from increasing model complexity.

Limitations

- Limited compute and time: hyperparameter search spaces were kept small and only a few configurations were explored per model family.
- several variants were not tested (e.g., deeper MLPs, alternative pooling strategies, different pretrained encoders).
- Limited number of training runs and random seeds, so results may be sensitive to initialization and specific splits.

References

- PolyAI, *BANKING77 Dataset*. Hugging Face Datasets, 2024. Available at: <https://huggingface.co/datasets/PolyAI/banking77>
- P. Schmid, *BERT-Banking77*. Hugging Face, 2021. Available at: <https://huggingface.co/philschmid/BERT-Banking77>
- N. Procopio, *DistilBERT-base-uncased-Banking77*. Hugging Face, 2023. Available at: <https://huggingface.co/nickprock/distilbert-base-uncased-banking77-classification>
- N. Procopio, *XLM-RoBERTa-base-Banking77*. Hugging Face, 2023. Available at: <https://huggingface.co/nickprock/xlm-roberta-base-banking77-classification>